How Many Visual Levers Drive Urban Perception? Interventional Attribution via Localised Counterfactual Edits

Jason Tang Stephen Law
University College London (UCL)
{jason.tang, stephen.law}@ucl.ac.uk

Abstract

Street-view perception models predict attributes such as safety, yet do not explain which scene elements should change to shift judgement. We propose interventional attribution*: a protocol that tests structured single-lever edits rather than correlational saliency over pixels or regions. Each lever specifies a semantic concept, grounded scene support, edit direction, and constrained prompt template. Edited candidates are retained only if they pass validity checks for same-place preservation, locality, realism, and plausibility. In a 50-scene pilot across five cities, 177 of 250 proposed edits pass this gate and 95 of those exceed an auxiliary safety threshold. Mobility Infrastructure and Physical Maintenance are the most promising families in the current setup. The contribution is the protocol and pilot benchmark, not a human-validated causal claim: pairwise human evaluation remains the ground-truth endpoint.*

1. Introduction

Street-view perception models predict subjective attributes such as safety, wealth, or liveliness at scale [3, 13, 17], yet they remain limited as explanatory tools: they do not answer the question central to urban design, namely *for this specific scene, which visual elements drive the perception, and what targeted modification would plausibly shift it?* Existing explainability methods such as saliency, SHAP, and concept probes remain correlational, identifying features associated with a score without validating that manipulating them changes human judgement [1, 7]. Empirical urban-design research links greenery, maintenance, and visibility cues to perceived safety [5, 6, 9, 14, 15], but does not by itself establish which localized visual changes validly shift perception for a specific scene.

Following Jacobs and defensible-space theory [5, 14], urban perception is often discussed as depending on multiple visible cues within the same scene. We test a restricted version of this idea: whether a given scene admits multiple distinct *single-lever* interventions that each, evaluated independently against the same unedited original, plausibly shift perception.

We propose an *interventional attribution* protocol that makes this question empirically testable. Rather than treating arbitrary regions as the explanatory unit, we define structured lever interventions drawn from an editor-feasible intervention vocabulary, each specifying a semantic concept, spatial support, intervention direction, and constrained edit template. Candidate edits are realised via prompt-conditioned image editing and retained only if a vision–language critic judges them to preserve the same place, remain localized, and appear realistic and plausible. Because generative editors may introduce correlated non-target modifications, each edit is treated as a hypothesis to be audited rather than as faithful evidence by default.

This paper does not claim validated causal effects of specific edits on human perception. Instead, it introduces an auditable interventional attribution framework and a pilot benchmark for studying that question. The current study uses a validity-gated generation pipeline and an auxiliary perception proxy to prioritize promising interventions for future human evaluation. Our empirical questions are: which lever families produce directional improvement on an auxiliary proxy once edits pass validity auditing, and across scenes, how many distinct single-lever interventions remain promising after validity screening and thresholding?

2. Related Work

Urban perception and design cues. Crowdsourced pairwise-judgement datasets have enabled vision models for perceived safety, wealth, and liveliness at scale [3, 13, 16, 17], but these models do not answer mechanistic scene-level questions. Urban-design theory and empirical street-view studies link perceived safety to surveillance, active edges, upkeep, greenery, and lighting [2, 5, 6, 9, 10, 14, 15], yet they stop short of testing localized, validity-audited single-lever interventions in a specific scene.

Interpretability and counterfactual editing. Saliency and

SHAP-style methods can be visually plausible yet unfaithful [1], and concept-based methods such as TCAV [7] remain correlational unless paired with explicit interventions. Counterfactual visual explanations [4], causal evaluation benchmarks [11], and urban counterfactual work [8, 12, 18] move closer to intervention, but typically optimize model outputs, feature-space perturbations, or full-scene generation. Our focus is narrower: identify which localized urban edits can be realised cleanly, preserved as the same place, and prioritized for later human evaluation. In that sense, the target is not a single best counterfactual, but a scene-grounded set of auditable lever interventions that can be compared and filtered before any human-endpoint study.

3. Method

Problem setup. Let x denote a street-view image and a a perceptual attribute such as safety. Our explanatory atom is a *lever intervention* $l = (c, s, d, \tau)$, where c is a lever concept, s the grounded scene support, d the intervention direction, and τ a constrained edit template. For image x, a planner grounds a bounded intervention vocabulary to the visible scene, yielding candidates $L(x)$. Each lever is tested independently against the unedited original; edits are never composed sequentially.

Phase 1: Lever candidate construction. Following the 9-page formulation, the intervention vocabulary is *editor-feasible*: each concept must satisfy three criteria, **(i)** empirical grounding in urban-perception literature [5, 6, 9, 14, 15], **(ii)** semantic distinctness from other concepts, and **(iii)** prompt-only editability without masks or geometry redesign. The current pilot uses twelve concepts in four families: Physical Maintenance, Environmental Amenity, Visual Legibility, and Mobility Infrastructure. The full vocabulary is listed in Appendix A.2.

Phase 2: Bounded stochastic intervention. Each candidate is edited independently from the original image with a bounded retry budget. A candidate enters the retained set

$$V(x) = \{l_i \in L(x) \mid \exists\, \tilde{x}_i \text{ passing the validity audit}\} \quad (1)$$

only if at least one edited sample passes four checks: same-place preservation, locality, realism, and plausibility. This validity gate is the key design choice. Generative editors can introduce non-target changes, so failed edits are excluded from attribution rather than being treated as noisy evidence. We report valid-edit coverage as $|V(x)|/|L(x)|$, but interpret it as a property of edit realizability under bounded prompt control rather than an urban-perception signal in itself.

Phase 3: Auxiliary scoring. For each retained edit $\tilde{x}_i$, we compute an auxiliary shift

$$\Delta_i = f_a(\tilde{x}_i) - f_a(x), \quad (2)$$

where f_a is a pretrained Place Pulse safety scorer. This is strictly an auxiliary proxy, not the main estimand, so Δ_i is used only as a prioritization signal. The main reported set is

$$E_{\text{aux}}(x, a) = \{l_i \in V(x) \mid \Delta_i > \theta_{\text{aux}}\}, \quad (3)$$

with $\theta_{\text{aux}} = 0.1$ in this pilot. The method therefore separates three questions: which levers can be grounded, which can be realised cleanly, and which retained edits look promising enough to send to human pairwise evaluation.[1]

4. Preliminary Results

Dataset and setup. We draw our image pool from SPECS (Street Perception Evaluation Considering Socioeconomics) [16], a demographically balanced perception survey in which 1,000 participants from five countries rated street-view scenes on ten perceptual indicators—six traditional (safe, lively, wealthy, beautiful, boring, depressing) and four new (live nearby, walk, cycle, green). SPECS provides per-image perception scores grounded in diverse human judgements, making it a principled source for studying which visual levers shift a specific attribute. We sample N=50 scenes from five SPECS cities (Amsterdam, Abuja, San Francisco, Santiago, Singapore; 10 per city), stratified jointly by baseline safety score and visual complexity so that the pilot spans the full range of starting conditions—from low-safety scenes where intervention is most needed to high-safety scenes that may offer fewer actionable levers. Each scene allows up to K=5 lever candidates with a stochastic budget of T=5 generation attempts per candidate. In this paper, we instantiate the framework with Qwen-Image-Edit as the prompt-only image editing model and GPT-5.4 as the LLM-as-judge critic to establish a reproducible lower-bound baseline. Validity auditing serves as an eligibility gate throughout this section: all effect summaries are computed only on edits retained in $V(x)$.

Auxiliary directional shifts. Among the 177 valid edits, the auxiliary scorer produces a mean shift of $+0.366$ (95% CI $[+0.199, +0.537]$, median $+0.184$, range $[-3.624, +3.842]$). Using threshold θ_{aux}=0.1, 95 of 177 valid edits fall in $E_{\text{aux}}(x, a)$, corresponding to a mean of 1.90 auxiliary-promising levers per scene (95% CI $[1.476, 2.324]$); 40 of the 50 scenes retain at least one θ_{aux}=0.1 auxiliary-promising lever and 24 retain multiple. In this pilot, the largest retained auxiliary shifts are observed for lane-marking repainting, crosswalk repainting, and localized greenery addition. At the same time, some valid edits remain neutral or negative, reinforcing that valid realization is necessary for analysis but not sufficient for positive auxiliary movement.

Table 1 shows that the main family-level signal lies in the Δ_{aux} column. Mobility Infrastructure yields the

[1] This is constrained interventional search, not causal identification or pixel-level decomposition. Distinct levers are tested independently against the same original image and are not assumed to add linearly.

Table 1. Lever-type results grouped by intervention family. Mean Δ_{aux} and Δ_{aux} [95% CI] are computed over valid edits only; confidence intervals are bootstrap percentile intervals over the retained set. "–" denotes concepts not proposed in this scene pool.

Family / Lever concept	Valid	Mean Δ_{aux}	Δ_{aux} [95% CI]
Physical Maintenance			
Graffiti removal	4	+0.396	[+0.121, +0.670]
Litter removal	25	+0.296	[−0.128, +0.743]
Facade repair	4	+0.519	[+0.105, +0.949]
Surface cleaning	38	+0.352	[+0.026, +0.684]
Family total	71	+0.344	[+0.107, +0.584]
Environmental Amenity			
Localized greenery add.	32	+0.650	[+0.241, +1.070]
Lighting repair	8	−0.232	[−0.598, +0.173]
Tree canopy management	14	−0.245	[−0.952, +0.446]
Family total	54	+0.287	[−0.041, +0.610]
Visual Legibility			
Signage decluttering	5	−0.631	[−1.294, −0.151]
Storefront transparency	2	+0.958	[+0.184, +1.732]
Family total	7	−0.177	[−0.927, +0.599]
Mobility Infrastructure			
Crosswalk repainting	15	+0.767	[+0.255, +1.261]
Lane marking repainting	30	+0.485	[+0.045, +0.934]
Family total	45	+0.579	[+0.244, +0.923]
Overall	177	+0.366	[+0.199, +0.537]

largest mean auxiliary shift in this pilot (+0.579, 95% CI [+0.244, +0.923]), with crosswalk repainting (+0.767, 95% CI [+0.255, +1.261]) and lane-marking repainting (+0.485, 95% CI [+0.045, +0.934]) leading the study. Physical Maintenance is the broadest consistently positive family (+0.344, 95% CI [+0.107, +0.584]), driven by surface cleaning, litter removal, facade repair, and graffiti removal. Environmental Amenity is more heterogeneous: localized greenery addition is strongly positive (+0.650, 95% CI [+0.241, +1.070]), whereas lighting repair and tree canopy management remain weak or negative on average. Visual Legibility remains small-sample and mixed in sign, so its directional pattern is not yet stable; the negative mean for signage decluttering (−0.631, 95% CI [−1.294, −0.151]) may reflect the auxiliary model associating commercial signage density with activity and perceived safety in Place Pulse training data. The Valid column is included to show the retained sample supporting each estimate; proposal and valid-rate diagnostics are reported in Appendix Tables A1 and A2, so the main table here remains focused on directional proxy movement.

Thresholded auxiliary-promising lever counts. Figure 1 is the main per-scene summary because it counts levers that are both valid and directionally promising under the auxiliary proxy at increasingly strict auxiliary thresholds. At the study threshold θ_{aux}=0.1, Mobility Infrastructure averages 0.56 auxiliary-promising levers per scene with mean retained shift +1.270 (95% CI [+0.958, +1.583]), Environmental Amenity averages 0.54 with mean +1.168 (95% CI [+0.821, +1.516]), and Physical Maintenance averages 0.76 with mean +1.044 (95% CI [+0.793, +1.296]). Visual Legibility averages only 0.04 levers per scene, so its interval

Table 2. City-level summary over the full N=50 run. Mean Δ_{aux} and Δ_{aux} [95% CI] are computed over valid edits only. Coverage diagnostics are reported in Appendix Table A2.

City	Valid	Mean Δ_{aux}	Δ_{aux} [95% CI]
Amsterdam	40	+0.400	[+0.073, +0.764]
Abuja	32	−0.136	[−0.441, +0.141]
San Francisco	36	+0.704	[+0.285, +1.105]
Santiago	38	+0.742	[+0.398, +1.089]
Singapore	31	−0.012	[−0.383, +0.352]
Overall	177	+0.366	[+0.199, +0.537]

is necessarily wide. By city, San Francisco and Santiago dominate the positive tail at θ_{aux}=0.1 with averages of 2.5 and 2.6 auxiliary-promising levers per scene respectively, while Abuja averages 1.1. As the cutoff rises to 1.0, the family ranking compresses but Environmental Amenity and Mobility Infrastructure still retain the largest positive tails in this pilot, and Amsterdam, San Francisco, and Santiago remain the cities with the highest high-delta leverage per scene.

Table 2 is best read as supporting context rather than a primary preliminary signal. Directional auxiliary movement is more context-dependent than validity gating: Santiago and San Francisco have the largest positive mean shifts, whereas Abuja and Singapore are near zero or slightly negative on average. The valid-rate differences remain informative as a feasibility diagnostic, with the appendix table showing Amsterdam and Santiago highest on retained valid counts and Singapore lowest.

Baseline score and gate-qualified editability. Appendix Figure A3 plots each scene's baseline safety score $f_a(x)$ against its valid lever count $|V(x)|$. At N=50, no meaningful monotonic relationship is observed (Spearman ρ=0.10, p=0.511). In other words, once generation reliability is no longer the binding constraint, scenes with lower baseline safety do not obviously admit more valid levers than higher-baseline scenes. We treat this as a secondary diagnostic about editability rather than a main perception result and revisit the null pattern in §5.

Validity gate. Of 250 proposed candidates, 177 pass the validity audit (rate 0.708); all 50 scenes retain at least one valid lever. Detailed coverage and failure analysis are in Appendix A.1.

Qualitative examples. Appendix Figure A2 shows three positive auxiliary-shift examples from distinct scenes: lane-marking repainting in Amsterdam, facade repair in Singapore, and surface cleaning in Santiago. All three are spatially constrained, visually legible, and plausible as single-lever urban changes.

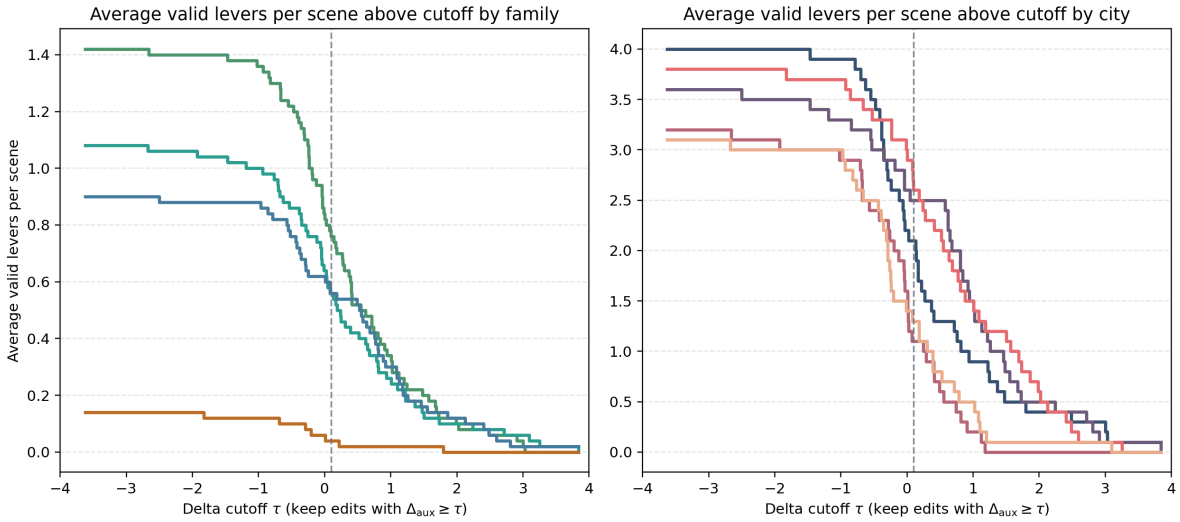

Figure 1. Average number of retained levers per scene satisfying $\Delta_{\text{aux}} \geq \tau$, split by lever family (left) and city (right). The annotation reports the overall mean Δ_{aux} with 95% CI and the mean number of θ_{aux}=0.1 auxiliary-promising levers per scene with 95% CI.

5. Discussion and Future Work

This paper establishes a pilot baseline with two outputs: a feasibility profile for the generation-and-audit pipeline, and a set of preliminary proxy-based directional patterns to prioritize for future human evaluation. The first is a generator-feasibility result: roadway markings and maintenance interventions realize reliably across cities, lighting repair remains the clearest bottleneck (valid rate 0.30), and the null baseline-score relationship (Spearman ρ=0.10) suggests that multi-lever richness is driven more by scene geometry and visible editable supports than by starting safety level. The second layer is more substantive: Mobility Infrastructure produces the largest retained auxiliary safety shifts under the current Place Pulse proxy. Whether that reflects genuine human perceptual sensitivity to roadway-legibility cues, scorer-specific priors, or an interaction with edit realizability remains open.

The heterogeneity within Environmental Amenity is equally informative. Localized greenery addition is strongly positive, while lighting repair and canopy management are weak or negative on average despite being theory-grounded cues. That pattern suggests family-level aggregation can hide important lever-specific differences in editability and scorer response, and motivates finer-grained prioritization at the lever level rather than the family level alone.

On the auxiliary proxy. The auxiliary scorer f_a has not been validated on diffusion-edited images, and Δ_{aux} should not be read as a calibrated perceptual effect size. The claim here is narrower: for semantically targeted, monotone interventions such as repainting markings, removing surface disorder, or adding greenery, the directional signal is grounded in a model trained on large-scale crowdsourced pairwise human preference over closely related scene variation, and the sign and relative magnitude of Δ_{aux} are useful for prioritizing which retained edits to send to human evaluation first.

XAI contribution. Beyond urban perception, the main contribution is methodological. The framework reframes visual explanation around structured, theory-grounded interventions rather than post hoc attribution over pixels or regions, and it imposes a stronger standard in which a proposed change must be semantically coherent and generatively feasible before entering the attribution result.

Future directions and human endpoint. The present pipeline is intentionally minimal: prompt-only generation, a prompted VLM critic, and an auxiliary proxy scorer establish a reproducible lower-bound baseline rather than a fully validated explanation system. The next steps follow directly from that structure: planner triage to allocate the candidate budget more intelligently, stronger generation control for validity, critic calibration against human judgements, pairwise human evaluation as the ground-truth endpoint, and eventually end-to-end optimization once planner, generator, and critic are each validated independently. Concretely, the 9-page analysis points to four near-term workstreams: predicting lever viability from scene features so the candidate budget is spent more selectively; improving generator alignment on known failure modes such as lighting repair; calibrating the critic against human E1 raters and 2AFC judgements; and only then considering joint optimization across planner, generator, and critic.

6. Conclusion

We presented a lever-based interventional attribution framework for street-view perception that reframes scene explain-

ability as a bounded search over structured, theory-grounded counterfactual edits. Rather than treating saliency over pixels or regions as the explanatory endpoint, the framework uses semantically coherent, validity-audited single-lever interventions as the unit of analysis. In a pilot over 50 scenes, this formulation yields a reproducible lower-bound baseline, a practical failure taxonomy, and preliminary proxy-based signals that can guide future human evaluation.

We view the present study as a pilot benchmark rather than a completed human-validated explanation system. Its value lies in formalizing the intervention unit, making the generation-and-audit process explicit, and identifying where future work should focus: planner triage, stronger editing control, critic calibration, and pairwise human judgement as the ground-truth endpoint. The substantive claim is deliberately narrow: once the validity gate is enforced, the pilot does reveal recurring directional patterns that help rank lever families, scenes, and cities for follow-on human study, even though those proxy shifts are not yet the human endpoint themselves.

References

[1] Julius Adebayo, Justin Gilmer, Michael Christoph Muelly, Ian Goodfellow, Moritz Hardt, and Been Kim. Sanity checks for saliency maps. In *Advances in Neural Information Processing Systems*, pages 9505–9515, 2018.

[2] Marco De Nadai, Radu-Laurentiu Vieriu, Gabriele Zen, Suzana Dragicevic, Nikhil Naik, Michele Caraviello, Cesar A. Hidalgo, Nicu Sebe, and Bruno Lepri. Are safer looking neighborhoods more lively? a multimodal investigation into urban life. In *Proceedings of the 24th ACM International Conference on Multimedia*, 2016.

[3] Abhimanyu Dubey, Nikhil Naik, Devi Parikh, Ramesh Raskar, and Cesar A. Hidalgo. Deep learning the city: Quantifying urban perception at a global scale. In *Computer Vision – ECCV 2016*, pages 196–212. Springer, 2016.

[4] Yash Goyal, Ziyan Wu, Jan Ernst, Dhruv Batra, Devi Parikh, and Stefan Lee. Counterfactual visual explanations. In *Proceedings of the 36th International Conference on Machine Learning*, pages 2376–2384, 2019.

[5] Jane Jacobs. *The Death and Life of Great American Cities*. Random House, 1961.

[6] Bin Jiang, Cecilia Nga Sze Mak, Hua Zhong, Linda Larsen, and Christopher John Webster. From broken windows to perceived routine activities: Examining impacts of environmental interventions on perceived safety of urban alleys. *Frontiers in Psychology*, 9:2450, 2018.

[7] Been Kim, Martin Wattenberg, Justin Gilmer, Carrie Cai, James Wexler, Fernanda Viegas, and Rory Sayres. Interpretability beyond feature attribution: Quantitative testing with concept activation vectors (tcav). In *Proceedings of the 35th International Conference on Machine Learning*, pages 2668–2677, 2018.

[8] Stephen Law, Rikuo Hasegawa, Brooks Paige, Chris Russell, and Andrew Elliott. Explaining holistic image regressors and classifiers in urban analytics with plausible counterfactuals. *International Journal of Geographical Information Science*, 37:2575–2596, 2023.

[9] Xiaojiang Li, Chuanrong Zhang, and Weidong Li. Does the visibility of greenery increase perceived safety in urban areas? evidence from the place pulse 1.0 dataset. *ISPRS International Journal of Geo-Information*, 4(3):1166–1183, 2015.

[10] Linda J. Loewen, G. Daniel Steel, and Peter Suedfeld. Perceived safety from crime in the urban environment. *Journal of Environmental Psychology*, 13(4):323–331, 1993.

[11] Thomas Melistas, Nikos Spyrou, Nefeli Gkouti, Pedro Sanchez, Athanasios Vlontzos, Yannis Panagakis, Giorgos Papanastasiou, and Sotirios A. Tsaftaris. Benchmarking counterfactual image generation. In *Advances in Neural Information Processing Systems, Datasets and Benchmarks Track*, 2024.

[12] Felipe A. Moreno, Andres De La Puente, and Jorge Poco. Urbanphysicaldisorder-4k: Understanding urban perception via counterfactuals and street view signs of physical disorder. In *IEEE International Conference on Big Data*, pages 5194–5200, 2025.

[13] Nikhil Naik, Jade Philipoom, Ramesh Raskar, and Cesar A. Hidalgo. Streetscore – predicting the perceived safety of one million streetscapes. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition Workshops*, pages 779–785, 2014.

[14] Oscar Newman. *Defensible Space: Crime Prevention through Urban Design*. Macmillan, New York, 1972.

[15] Boris A. Portnov, Rasha Saad, Tal Trop, Doron Kliger, and Anna Svechkina. Linking nighttime outdoor lighting attributes to pedestrians' feeling of safety: An interactive survey approach. *PLOS ONE*, 15(11):e0242172, 2020.

[16] Matias Quintana, Youlong Gu, Xiucheng Liang, Yujun Hou, Koichi Ito, Yihan Zhu, Mahmoud Abdelrahman, and Filip Biljecki. Global urban visual perception varies across demographics and personalities. *Nature Cities*, 2(11):1092–1106, 2025.

[17] Philip Salesses, Katja Schechtner, and Cesar A. Hidalgo. The collaborative image of the city: Mapping the inequality of urban perception. *PLOS ONE*, 8(7):e68400, 2013.

[18] Chenbo Zhao, Yoshiki Ogawa, Shenglong Chen, Takuya Oki, and Yoshihide Sekimoto. Street space quality improvement: Fusion of subjective perception in street view image generation. *Information Fusion*, 125:103467, 2026.

Table A1. Family-level summary over the full $N{=}50$ run. Valid-rate intervals are 95% Wilson confidence intervals; mean Δ_{aux} and Δ_{aux} [95% CI] are computed over valid edits only.

Family	Valid / Prop.	Rate [95% CI]	Mean Δ_{aux}	Δ_{aux} [95% CI]
Physical Maintenance	71 / 92	0.77 [0.68, 0.85]	$+0.344$	$[+0.107, +0.584]$
Environmental Amenity	54 / 82	0.66 [0.55, 0.75]	$+0.287$	$[-0.041, +0.610]$
Visual Legibility	7 / 10	0.70 [0.40, 0.89]	-0.177	$[-0.927, +0.599]$
Mobility Infrastructure	45 / 66	0.68 [0.56, 0.78]	$+0.579$	$[+0.244, +0.923]$
Overall	177 / 250	0.71 [0.65, 0.76]	$+0.366$	$[+0.199, +0.537]$

Table A2. City-level coverage and directional summary over the full $N{=}50$ run. Each city contributes 10 scenes and therefore 50 proposed candidates. Valid-rate intervals are 95% Wilson confidence intervals; mean Δ_{aux} and Δ_{aux} [95% CI] are computed over valid edits only.

City	Valid / Prop.	Rate [95% CI]	Mean Δ_{aux}	Δ_{aux} [95% CI]
Amsterdam	40 / 50	0.80 [0.67, 0.89]	$+0.400$	$[+0.073, +0.764]$
Abuja	32 / 50	0.64 [0.50, 0.76]	-0.136	$[-0.441, +0.141]$
San Francisco	36 / 50	0.72 [0.58, 0.83]	$+0.704$	$[+0.285, +1.105]$
Santiago	38 / 50	0.76 [0.63, 0.86]	$+0.742$	$[+0.398, +1.089]$
Singapore	31 / 50	0.62 [0.48, 0.74]	-0.012	$[-0.383, +0.352]$
Overall	177 / 250	0.71 [0.65, 0.76]	$+0.366$	$[+0.199, +0.537]$

A. Appendix

A.1. Validity and Coverage Details

Across 50 scenes the planner produces the full 250 candidate rows (5.0 per image on average), of which 177 pass the validity audit, yielding mean coverage $|V(x)|/|L(x)| = 0.708$. Figure A1 shows the distribution of valid counts per scene: 2 scenes have one valid lever, 3 have two, 16 have three, 24 have four, and 5 have all five.

The dominant bottleneck is not candidate proposal but valid realisation. Among the 73 audited-but-rejected edits, plausibility dominates completely; 45 also show no discernible target change, 25 exhibit non-local drift, 5 are unrealistic, and none fail same-place preservation. Generator failure before audit is zero in the cleaned $N{=}50$ run, so the remaining variance is concentrated in valid realisation rather than missing candidate production.

A.2. Intervention Vocabulary

Each concept below satisfies all three inclusion criteria defined in §3. Structural interventions (building setback, road geometry, facade cladding) and social-order cues (pedestrian density, active frontage) are excluded because they are not realisable under prompt-only editing without masking or inpainting.

Physical Maintenance (broken-windows and perceived-upkeep literature [6, 14]): graffiti removal; litter removal; facade repair; surface cleaning; shutter repair.

Environmental Amenity (biophilic safety and nighttime visibility [9, 15]): localised greenery addition; lighting repair; tree canopy management.

Visual Legibility (active-edges and surveillance theory [5]): signage decluttering; storefront transparency increase.

Mobility Infrastructure (pedestrian safety and walkability [16]): crosswalk repainting; lane marking repainting.

A.3. Prompt Contracts

The planner, editor, and critic each operate under a structured prompt contract. Condensed versions are shown below; full prompts are in the released repository.

Planner prompt (condensed)

```
You are an urban perception planner.  Given a
street-view image and target percept, propose a
constrained set of candidate lever interventions.

ONTOLOGY: Choose only from four families ---
Physical Maintenance, Environmental Amenity, Visual
Legibility, Mobility Infrastructure --- and prefer
cross-family diversity when the scene supports it.

HARD CONSTRAINTS: (1) One lever per candidate.  (2)
Grounded in a visible scene element.  (3) Local,
plausible, prompt-only friendly.  (4) No global
relighting, weather, or camera changes.  (5) Prefer
the smallest plausible intervention.  (6) Exclude
theoretically relevant levers whose target element is
not clearly visible/editable in the image.

DIVERSITY: Return distinct candidates using different
lever concepts and avoid magnitude variants of the
same intervention.

Return JSON with field ``candidates'', each
containing:  lever_concept, lever_family,
scene_support, target_object, intervention_direction,
edit_template, edit_plan.
```

Edit prompt (condensed)

```
Use the PROVIDED image as base.  Preserve exact
viewpoint, geometry, and layout.

ALLOWED: Only modify the target object as required by
the plan.  If repainting/retexturing, keep shape and
placement identical.

FORBIDDEN: No global restyling, relighting, or
recoloring.  No adding/removing other objects.  No
readable text.  No background, sky, road, or context
changes.

Lever concept:  {lever_concept}
Lever family:  {lever_family}
Scene support:  {scene_support}
Intervention direction:  {intervention_direction}
Edit template:  {edit_template}
Target object:  {target_object}
Edit plan:  {edit_plan}
```

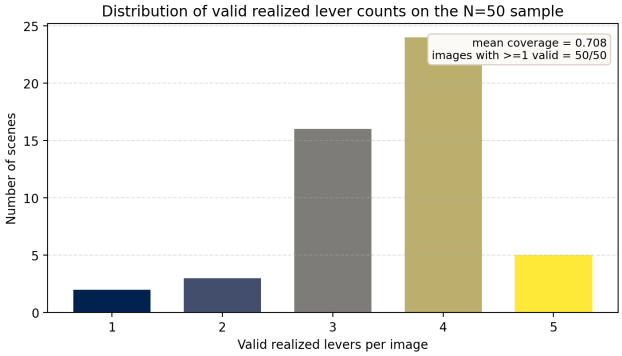

Figure A1. Distribution of valid realised lever counts per scene (N=50). Most scenes admit three or four independently valid single-lever interventions.

Critic prompt (condensed)

```
Evaluate whether the edited image is a valid
single-lever counterfactual relative to the original.

CONTEXT: Edit produced by prompt-only diffusion
without masking.  Minor incidental changes (tone
drift, texture resampling) are expected artefacts
and should not cause failure unless they materially
alter scene meaning; plausibility and locality are
still judged strictly against the requested lever and
support.

CRITERIA: (A) edit_attempted:  generator made
a visible change at the target (false only if
output looks identical to original there).  (B)
same_place_preserved:  same underlying place.
(C) is_localised:  primary meaningful change is
in/near intended support; minor global tone shifts
do not count as non-local.  (D) is_realistic:
physically plausible and coherent.  (E) is_plausible:
recognisably the requested lever at the stated
support; fail if the edit type/support is wrong or
too excessive.

CLEAR FAIL CONDITIONS: Viewpoint/geometry changed;
large non-target objects added/removed; requested
lever replaced by a different change; no discernible
change at target.

Return JSON: {edit_attempted, same_place_preserved,
is_localised, is_realistic, is_plausible,
notes}, where notes = {failure_modes, diagnosis,
repair_suggestion}.
```

Figure A2. Retained qualitative examples from the pilot. From top: lane-marking repainting (Amsterdam), facade repair (Singapore), and surface cleaning (Santiago).

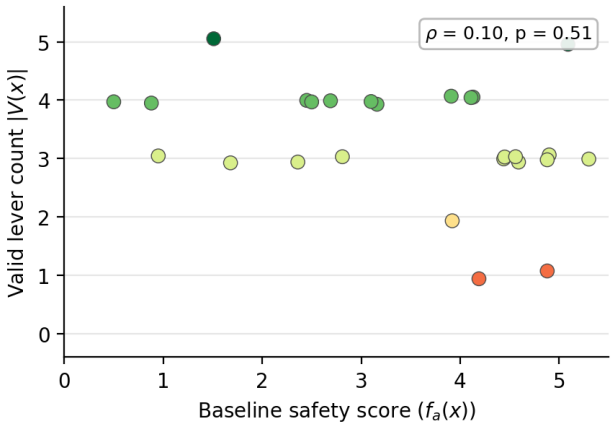

Figure A3. Baseline safety score vs. valid lever count per scene (N=50). No meaningful relationship is observed (Spearman ρ=0.10, p=0.511).